

Three things that shipped today.

01 A cache-reuse trick cuts long-context serving cost, and it is a config flag.

A cache-reuse trick cuts long-context serving cost, and it is a config flag.

- every request recomputes the shared prefix
- + the prefix is cached once, reused by all requests

WHY IT MATTERS

Serving cost stops scaling with total tokens and starts scaling with unique tokens.

WHO SHOULD CARE

Anyone serving one long system prompt, RAG preamble, or few-shot block to many users.

Nine lines. Runs fully offline, no model download.

```
def approx_tokens(text):  
    return max(1, round(len(text) / 4))  
  
prefix = "You are a helpful assistant. " * 40  
requests = [prefix + q for q in ("hi", "why is the sky blue?", "can you say that again?")]  
  
naive = sum(approx_tokens(r) for r in requests)  
shared = approx_tokens(prefix) + sum(approx_tokens(r[len(prefix):]) for r in requests)  
  
print(naive, shared, round(1 - shared / naive, 3))
```

881 302 0.657

65.7%

fewer tokens billed in the example below

65% of the tokens in that workload were the same prefix, paid for three times. The fix is a serving-side flag, not a new model.

[**FULL CODE IN DESCRIPTION**](#)[**A NEW DIFF EVERY WEEKDAY**](#)